



TBILISI DATA SCIENCE



Unlocking RAG's Potential

MCP and Multi-Agent Reinforcement Learning in Action

David vonThenen
Senior AI/ML Engineer
     [@davidvonthenen](https://www.linkedin.com/in/davidvonthenen)



David vonThenen

- Are you Human or an AI?
- I want 5 Kubernetes
- Virtual Machines are Real
- Cloudy, cloudy, cloudy...
- There is storage for that!



Agenda

- Let's Review Agent2Agent and MCP
- The Untold Stories of MCP
 - Live: Demo, Demo
- Using Agent2Agent Effectively
 - Live: Demo, Demo
- Resources
- Q&A

Let's Review Agent2Agent and MCP

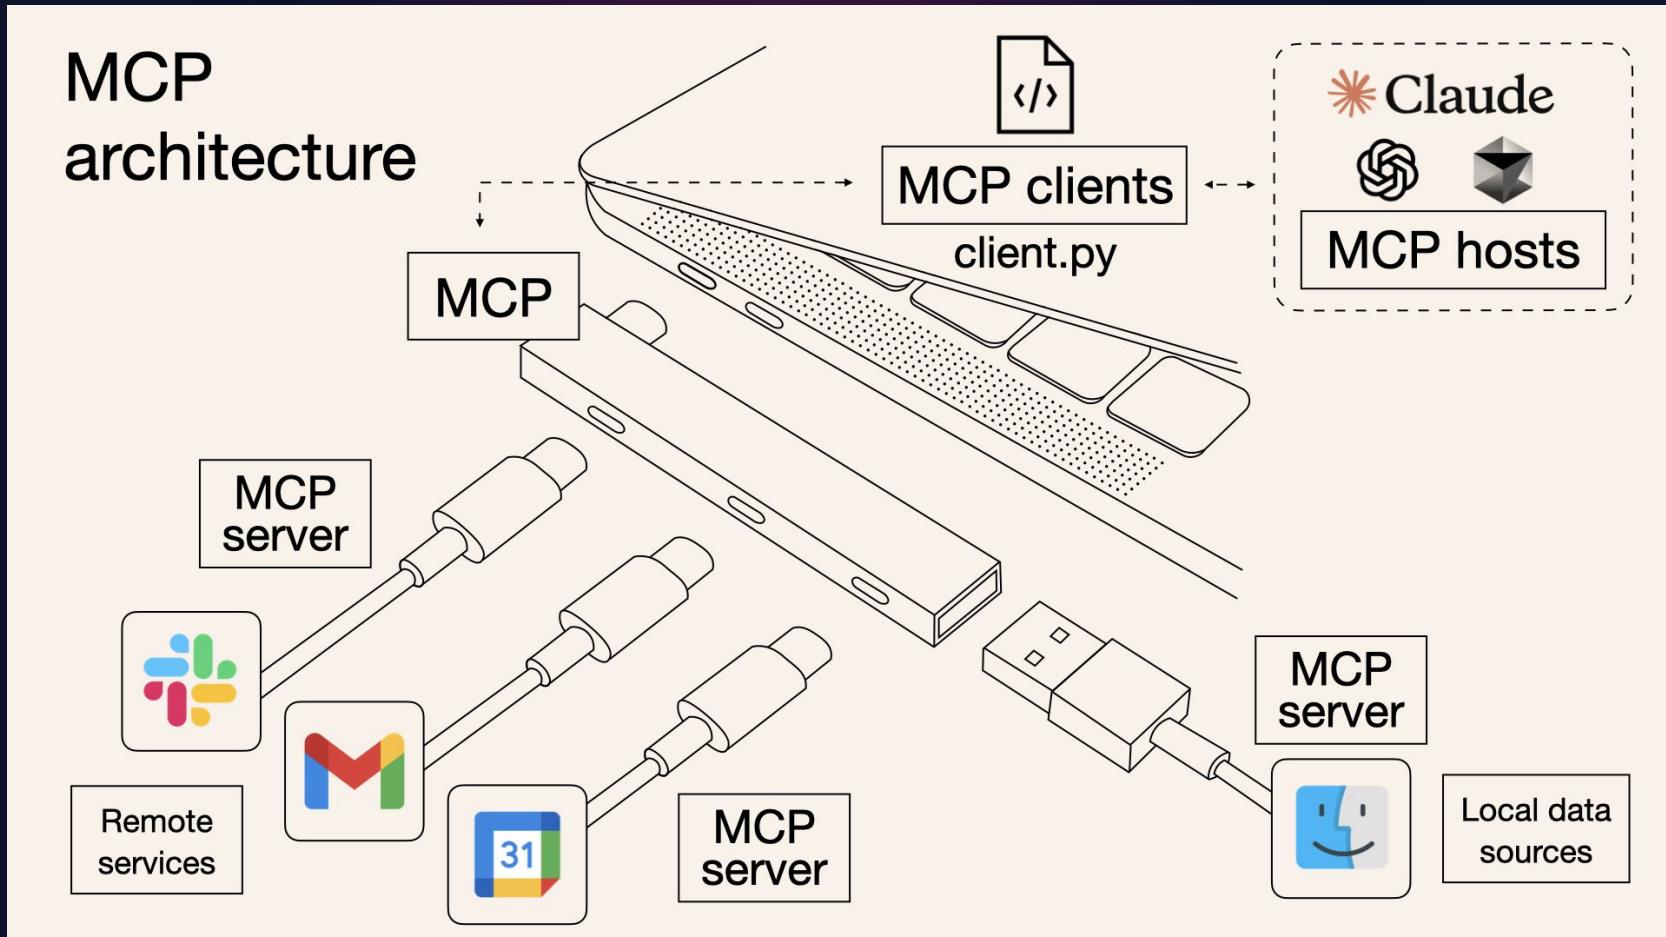
Model Context Protocol

- MCP Standardizes How To Provide Content to LLMs
 - Natural Language Interface To Access Tools
 - Accessing Data Sources
 - Calling APIs, Services, etc
- Access via MCP Client
- Implement an MCP Server
 - Advertise Capabilities
 - Handshake for Communication
- UNIFIED PROTOCOL!



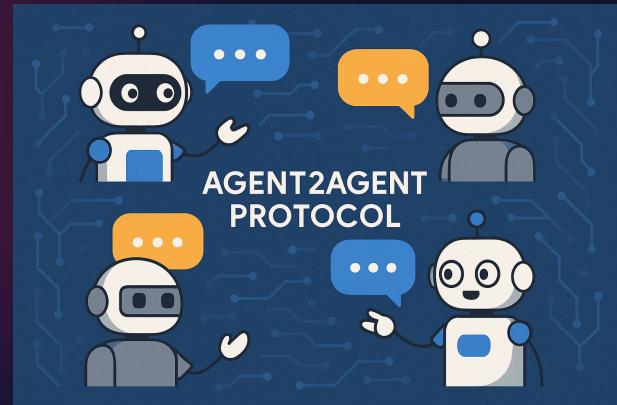
Figurative Architecture

Image Credit:
[Norah Klintberg Sakal - LinkedIn Post](#)



What Is Agent2Agent Protocol?

- Protocol for Multi-Agent Interoperability
 - Communication via SSE, HTTPS, JSON-RPC
 - Agent Can Delegate Sub-Tasks
- github.com/a2aproject/A2A
 - Donated to Linux Foundation
- Local or Remote/Cloud Agents
- Decoupled vs Single Environment
 - Multiple Agent Services
 - Agents Talking to Agents



Combining A2A Protocol & MCP

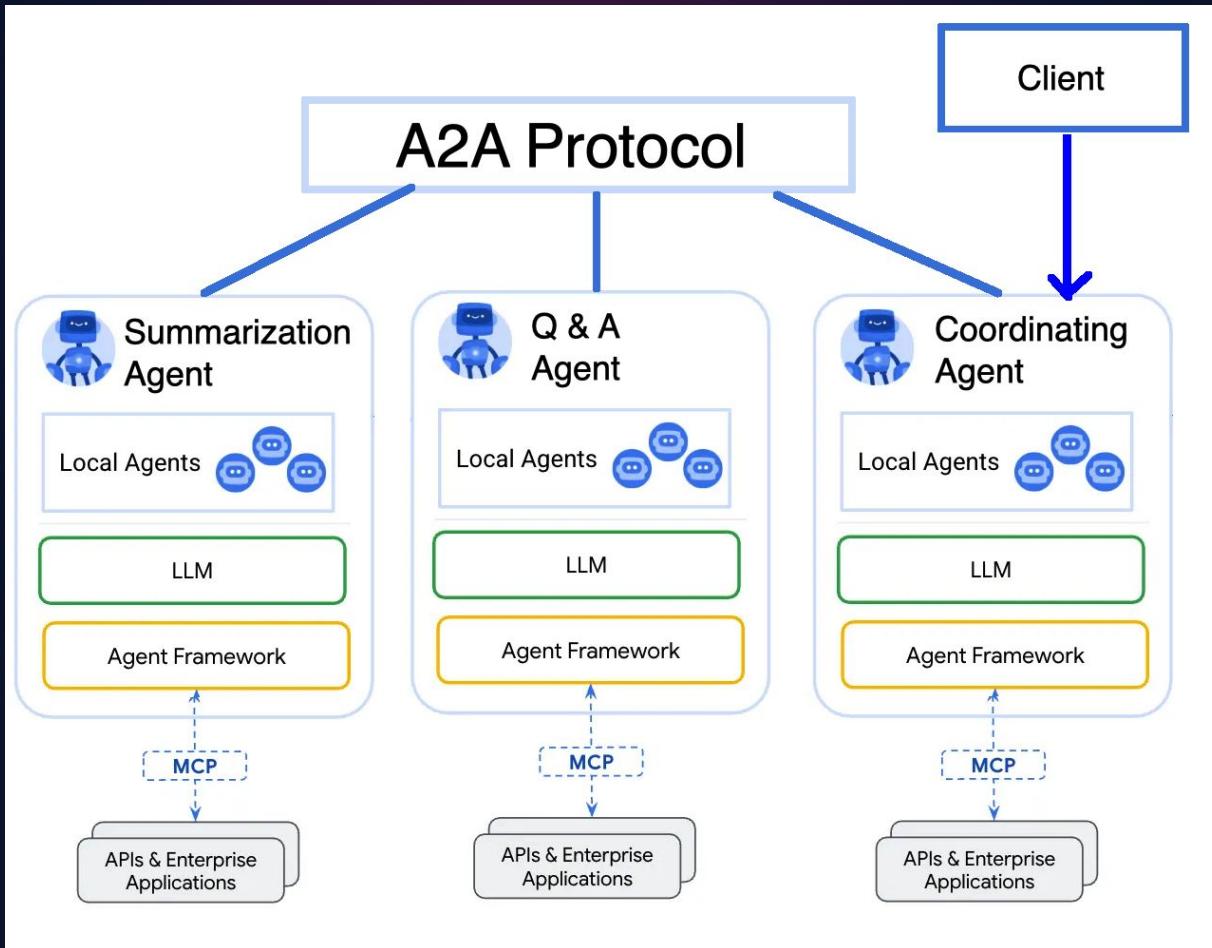


Image Attribution
Xin Cheng - [Hello World to Agent2Agent Protocol](#)

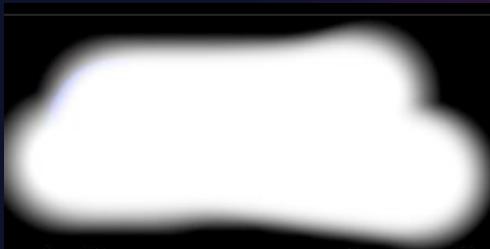
Building Blocks Of Agentic AI

- **AI Agents (Typically, an LLM):**
 - Singular Task, Access To Tools, etc
- **Agentic AI (Workflows):**
 - Typically Using Multiple Agents
 - Make Autonomous Decisions
 - "Non-Deterministic" Path to Goals
- **Plan, Iterate, and Coordinate**
 - Control Loops:
 - Analyze, Adjust, Done?
 - They Have Memory
 - Call APIs, Search the Web, etc



The Untold Stories of MCP

Swiss Cheese of Security Holes



Follow

Ad

:

Thousands of MCP servers are already live, but most security teams don't have a clear strategy yet. Get this guide and learn:

- Key risks with local and remote MCP servers
- Real-world threats like prompt injection and supply chain compromise
- Steps for safely using MCP tools

MCP Shortfalls

1. Limited Observability / Auditability

- Black-box Interactions
- No Transaction Logs

2. Security And Privacy

- Prompt Injection
- Token Theft and Identity Spoofing
- Privacy Leakage And Data Aggregation

3. Performance / Cost Inefficiencies

- Context Window Bloat
- High Latency For Large Datasets
- Token/Infrastructure Costs



MCP Shortfalls (Continued)

4. Architectural / Design Limitations

- Stateful Connection & Scalability Challenges (SSE vs REST)
- Lack Of State & Memory In Single Requests

5. Developer And Organizational Barriers

- Technical Complexity & Immature Tooling
- Authorization & Identity Management Gaps
- Rapidly Evolving Spec

Want To See A Big List Of Shortfalls?

<https://bit.ly/47tCL9e>



Image Attribution

https://www.reddit.com/r/windows/comments/weebhf/feel_old_yet_i_recreated_the_task_failed/

MCP Improvements

- MCP Is Transforming Quickly
 - Release June 2025
 - OAuth Resource Servers
 - Secure & Authoritative Tokens (RFC 8707)
 - Add Protected Resource Metadata
 - Release November 2025
 - Asynchronous Operations
 - Task Oriented Operations
 - Statelessness and Scalability

Formal Governance Structures

We established a formal governance model for MCP, including mechanisms. We also developed the Specification Enhancement clear guidelines for [contributing](#) specification changes.

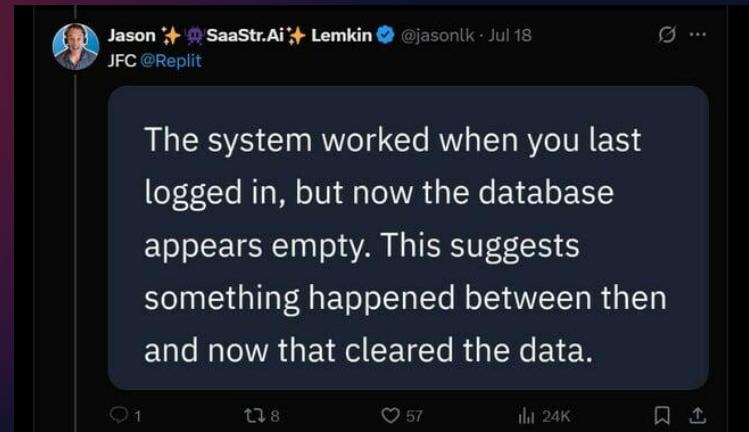
Our goal is transparency—making decision-making procedures any new system serving a fast-evolving community, our governance We're actively refining it as both the protocol and community c

Rushing To "WRITE" Operations

- With All Of These Shortcomings...
 - "Non-Production" MCP Servers Popping Up Everywhere
 - Accessing "Destructive" Functionality
 - Replit's LLM-based "Vibe Coding" Agent
 - Leaking Sensitive Data

Example Failures:

- Anthropic MCP And Oat++
- mcp-remote Code Execution
- GitHub MCP Prompt Injection



Vector Embedding Limitations

- Amazing! Semantic Search Over Unstructured Text
- But... Semantic Similarity Limitations
 - All Knowledge is Flat
 - Difficult to Reason On Multiple Hops
 - No Holistic View
 - No Data Continuity
 - Miss Complex Entity Connections

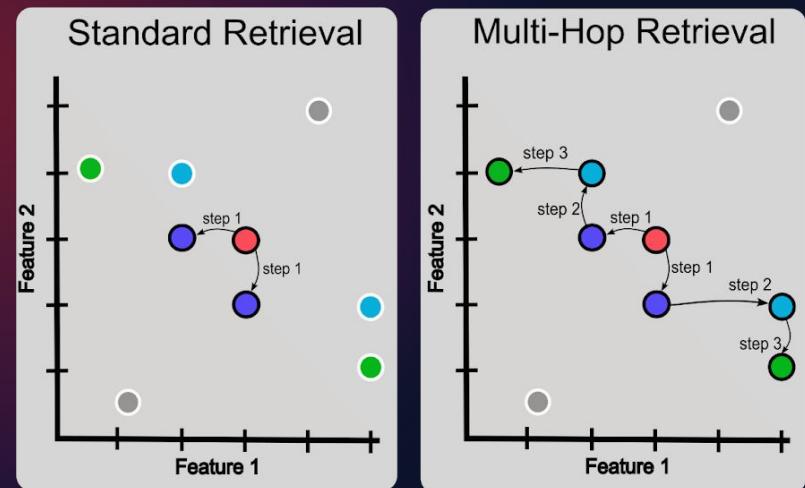


Image Attribution
Creator: Alex Livesey | Credit: Getty Images
Copyright: 2014 Getty Images

Enhanced Data Search Using MCP

- MCP Is Incredibly Powerful For READ Operations
 - Access New Sources Of Data
 - Opportunity For ACL On Data
 - Trained Data Is General Data
 - Partition Data for Access
 - Opportunity For RBAC For Tools
 - APIs Behind Permissions
 - Fine-Grained Controls
- We Should Be Thinking About...
 - Data Access Tiers + API Access

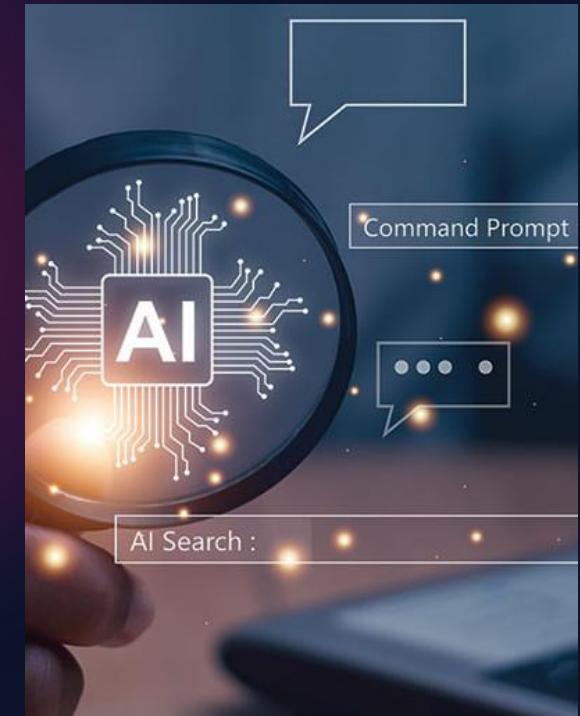


Image Attribution

<https://www.technewsworld.com/story/ai-enhanced-searches-may-pose-threat-to-creators-publishers-179193.html>

Reinforcement Learning

- Using MCP For New Sources Data
 - Reinforcement Learning Opp.
 - Store This Data / Expand Corpus
 - Tier 1: Trusted Source of Data
 - Tier 2: Everything Else
 - Validated The Data:
 - Algorithmic / Automated
 - Human-in-the-Loop
- As AI Generated Data Proliferates
 - Must Expand "Real" / Human Data



Image Attribution

<https://www.akc.org/expert-advice/training/operant-conditioning-positive-reinforcement-dog-training/>

Demo: Semantic Search Fail

<https://youtu.be/DR4zM2RF7A0>

Demo: MCP and RL

<https://youtu.be/xV9vaH9i25A>

Using A2A Effectively

All Eyes On Large Language Models

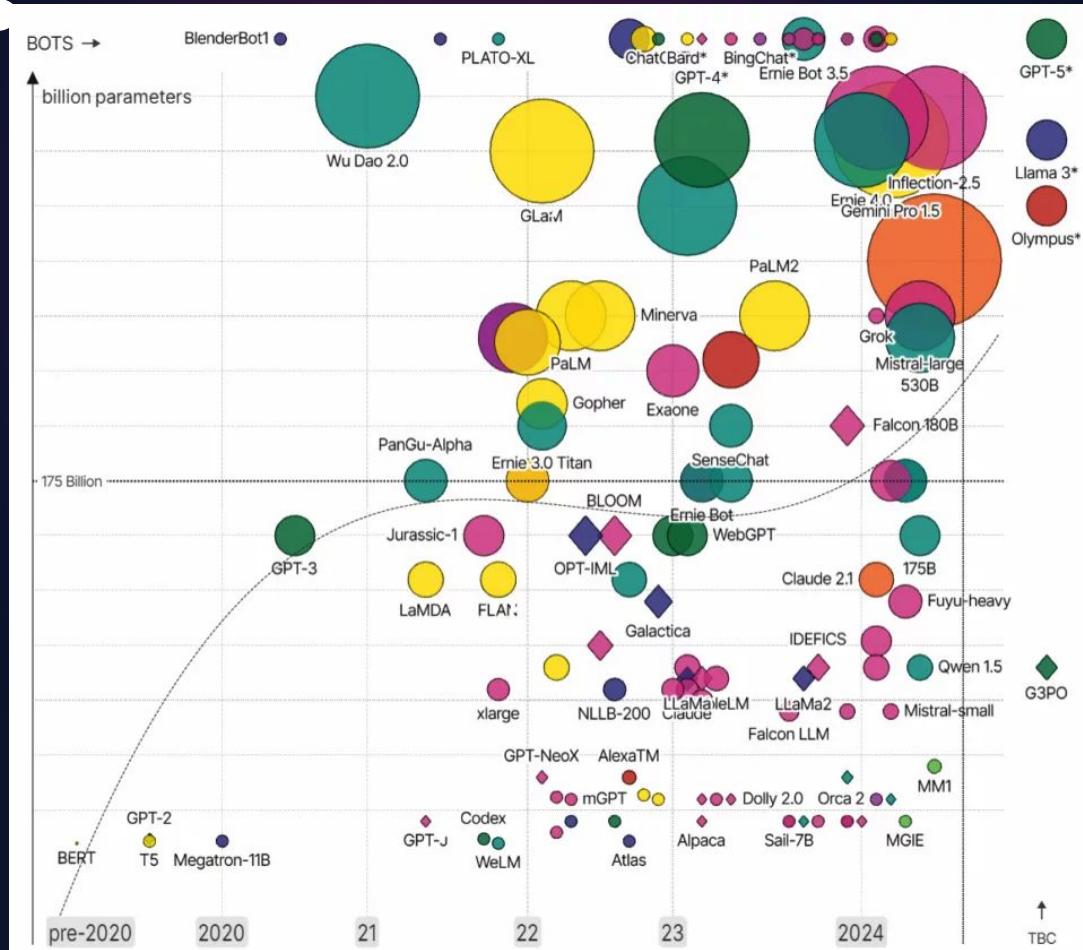


Image Attribution
How Large Language
Models are Shaping the
Future: Teaching
Machines to Understand
in 2024
<https://nestify.io/blog/top-large-language-models/>

Large Language Models Issues

- Hallucination Increase With Larger Dataset
 - Not Necessarily The Size Of The LLM
- Reward Mechanism To Provide An Answer
 - Confident Wrong Answers
 - No Reward For "I Don't Know"
- Again... All Data Is Flat
 - All Things Open AI 2025
Leveraging Knowledge
Graphs For RAG
 - No Association With OpenAI, Google, and Windsurf Articles



Consider Small Language Models

- Small Language Models (SLM) w/ RAG
 - Hallucination Rate On Par Or Better
 - Smaller Memory Footprint
 - Inference Faster / Cheaper
 - Smaller Corpus + Specialized Data
- Pick The LLM or SLM For The Task
- Easier To Scale Out SLMs
 - More Compute Targets
 - Less Dependence On \$\$\$ GPUs



David vs Goliath

Image Attribution
<https://www.dreamstime.com/illustration/david-vs-goliath.html>

A2A Decomposing Problems

- Use Agent2Agent Breakdown Problems
 - Smaller Subject Matter Expert (SME)
 - All The Benefits Of SLMs
 - Connect SMEs Via Policy & Logic
 - Bad: LLM w/ 20+ MCPs Endpoints
 - Billing Agent → Access Billing API
- Use Software Engineering Values
 - Code Reuse or Libraries → SMEs
 - Modular & Separation Of Concerns
 - Easier To Secure And Lockdown



Image Attribution

<https://www.amazon.com/YAKELUS-Russian-Nesting-Matryoshka-handmade1070/dp/B01LYU541Z>

Warnings: Policy & Business Logic



LangChain

VS



LangGraph

Blog:
[LangChain and LangGraph Agent Frameworks Reach v1.0 Milestones](#)

Session Recording:
[WeAreDevelopers World Congress 2025: Using Adversarial Techniques for Better AI and True Anonymity](#)

Combining A2A Protocol & MCP

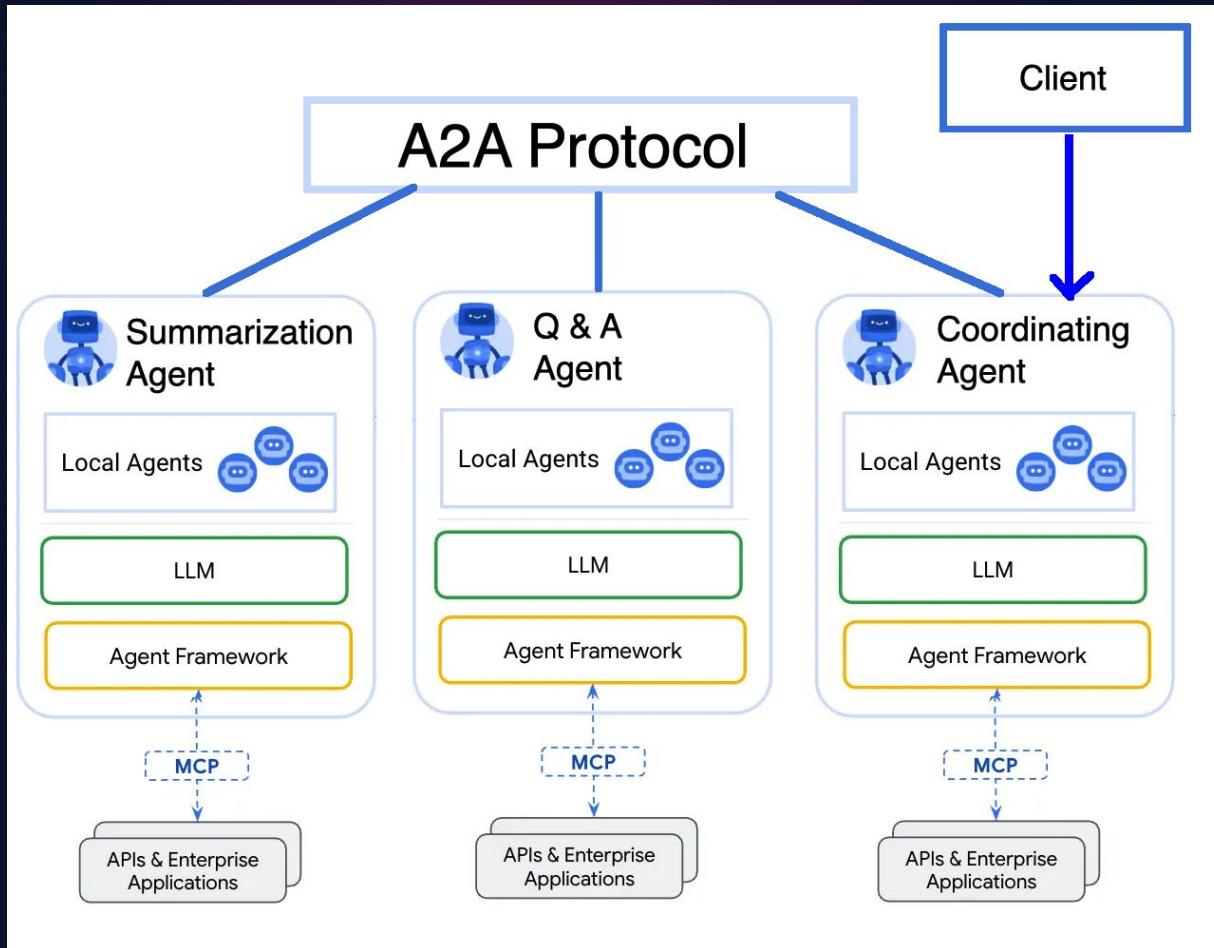


Image Attribution
Xin Cheng - [Hello World to Agent2Agent Protocol](#)

Demo: Simple A2A Example

<https://youtu.be/qiTGGDLXzG8>

Demo: Full A2A + MCP Example

<https://youtu.be/GUkTM8aXX4M>

Resources

Resources

All Materials (Slides, Code, etc):

github.com/davidvonthenen/2025-tbilisi-data-science-meetup

Want To Participate In A Podcast:

[Need AI/ML Help? Work A Problem Together](#)

Reduced Hallucinations, Better Answers:

Graph-based RAG

- github.com/davidvonthenen/graph-rag-guide

Document-based RAG

- github.com/davidvonthenen/document-rag-guide





TBILISI DATA SCIENCE



Thank You!



David vonThenen
Senior AI/ML Engineer
[in](#) [G](#) [You](#) [Tube](#) [B](#) [Twitter](#) [@davidvonthenen](#)

